

Data: Splits, Bias, Variance, ...

Denis Baskan
Zia Muhammad
Rashik Islam
Alexei Figueroa



BEUTH HOCHSCHULE
FÜR TECHNIK
BERLIN

University of Applied Sciences




Content

- **Data**
 - Amount
 - Splits
 - Distribution
- **Errors**
 - Irreducible Error
 - Bias and Variance Error
 - Underfitting and overfitting
 - Avoidable Bias
- **BV- Tradeoff**
 - In terms of MSE
 - Impact
 - Techniques to control
- **BV - Ensemble Methods and NNs**
 - From trees to forests
 - Neural networks


Data - Amount

- **ML & DL Methods are taking off now**




Data - Amount

Annual size of global dataspace




1 ZB = 1e9TB = 1,000,000,000TB


Data - Amount



Data - Amount






Data - Amount




Data - Amount

- Curse of Dimensionality




Data - Amount

- **Curse of Dimensionality**



Data - Amount

Big-O Complexity Chart



Data - Amount

| Algorithm | Classification/Regression | Training | Prediction |
|--------------------------------------|---------------------------|---------------------------|--|
| Decision Tree | C+R | $O(n^2p)$ | $O(p)$ |
| Random Forest | C+R | $O(n^2pn_{trees})$ | $O(pn_{trees})$ |
| Random Forest | R Breiman implementation | $O(n^2pn_{trees})$ | $O(pn_{trees})$ |
| Random Forest | C Breiman implementation | $O(n^2\sqrt{p}n_{trees})$ | $O(pn_{trees})$ |
| Extremely Random Trees | C+R | $O(npn_{trees})$ | $O(npn_{trees})$ |
| Gradient Boosting (n_{trees}) | C+R | $O(npn_{trees})$ | $O(pn_{trees})$ |
| Linear Regression | R | $O(p^2n + p^3)$ | $O(p)$ |
| SVM (Kernel) | C+R | $O(n^2p + n^3)$ | $O(n_{sv}p)$ |
| k-Nearest Neighbours (naive) | C+R | — | $O(np)$ |
| Nearest centroid | C | $O(np)$ | $O(p)$ |
| Neural Network | C+R | ? | $O(pn_{l_1} + n_{l_1}n_{l_2} + \dots)$ |
| Naive Bayes | C | $O(np)$ | $O(p)$ |

Splits


- **3 Data Sets**
 - **Training Set**
 - Used for training model
 - **Validation Set**
 - Subset of training set
 - Used for optimizing (tuning parameters, selecting features,...)
 - **Test Set**
 - Evaluate your model
- **Split Ratio 70/30 or 80/20**
 - Shuffle data & stratify by classes

Splits

- **Validation & Test Set**
 - Data that you expect in future
- **Example 1 - Turnover for next year**
 - Last 3-5 years might be helpful
- **Example 2 - Image Classifier**
 - User feeds images from smartphone
 - Model was trained on images from web

Data Distribution

- **Good Performance on Training/Validation Set**
- **Reasons for poor performance on test set**
 - **Overfit to validation set**
 - Use more data
 - Train less



Data Distribution

- **Good Performance on Training/Validation Set**
- **Reasons for poor performance on test set**
 - **Overfit to validation set**
 - Use more data
 - Train less
 - **Test set is more complex**
 - Model hasn't seen all the features
 - Model is too small / simple
 - **Test set comes from different distribution**
 - Example: Image is horizontally flipped



kfold cv?

<https://towardsdatascience.com/train-test-split-and-cross-validation-in-python-80b61beca4b6>

Data Distribution

- **When to use different distribution?**
 - Not enough data from customer / user
 - Some features might not be in data
 - Collect data from other sources
 - Split user data into 2 parts
 - Training / Validation
 - Testing

Data Distribution

- **Why not using all (user) data?**
 - Was the way to go in past.
 - Other data sources could bias your data
 - Other data sources could harm your model
 - Some features might not be in data
 - Collect data from other sources
 - Split user data into 2 parts
 - Training / Validation
 - Testing

Data Distribution

- **Using customer data only**
 - Was the way to go in past.
 - Other data sources could bias your data
 - Other data sources could harm your model
 - Training models takes more time
 - Adding more data
 - More features
 - Small model might not capture all features
 - Inconsistent data
 - Image classifier
 - Predicting house prices

Data Augmentation

Example: Image-based methods

- Horizontal flip
- Rotate
- Scale
- Crop
- Translation
- Noise
- Blur
- Brightness / Contrast
- RGB -> Gray -> BW


Basic Error Analysis

- **Irreducible Error**
- **Reducible Error**
 - **Bias Error**
 - **Variance Error**

What is Irreducible error?

Irreducible Error:

- The Measure of the amount of noise in the Data
- It is usually caused by unknown variables that may be having an influence on the output variable.
- How good we make our model, our data will have certain amount of noise or irreducible error that can not be removed.



What is Bias?

Bias:

The difference between the average prediction of our model and the correct value which we are trying to predict on training data.


- Poorly Perform on Training data
- Low Training Accuracy
- Example:
 - Training error = 15%
 - Test error = 16%

The bias as 15%, and variance as 1% (**Variance = Test error - Training error**) This classifier is fitting the training set poorly with 15% error, but it's error on the Test set is barely higher than the training error. This classifier therefore has **high bias**, but low variance.

Underfitting

Underfitting:

- Model can not capture underlying pattern of the data
- High Bias leads to Underfitting



Underfitted

Techniques To Reduce High Bias

Techniques To Reduce High Bias:

- Train Longer
- Train a more complex model
- Decrease Regularization
- New model architecture

Avoidable Bias

Avoidable Bias:

The difference between the training set error and the optimal error rate.

- The “avoidable bias” reflects how much worse your algorithm performs on the training set than the “optimal model.”
- Optimal error rate smallest possible error that the algorithm can reach.
- Difference (Training Error, Human-Level Performance) = Avoidable Bias
- Difference (Validation Error, Training Error) = Variance

Example: Classification Cat vs Not Cat

| | Classification error (%) | |
|-------------------|--------------------------|------------|
| | Scenario A | Scenario B |
| Humans | 1 | 7.5 |
| Training error | 8 | 8 |
| Development error | 10 | 10 |

| | Scenario A | Scenario B |
|----------------|------------|------------|
| Human Error | 1% | 7.5% |
| Avoidable Bias | 7% | 0.5% |
| Variance | 2% | 2% |

Techniques for avoidable Bias

Techniques to reduce avoidable Bias

- Increase the model size
- Reduce regularization
- Modify model architecture


What is Variance?

Variance:

Variance is the variability of Model Prediction for a given data point.

- Low Testing Accuracy
- Example
 - Training error = 1%
 - Test error = 11%
 - Variance = Test error - Training error

The bias as 1%, and the variance as 10% (=11% - 1%). Thus, it has **high variance**. The classifier has very low training error, but it is failing to generalize to the Test set.



Overfitting

Overfitting:

- Model capture underlying pattern too well of the training data.
- High Variance leads to Overfitting



Overfitted

How To Reduce High Variance?

High variance is due to a model that tries to fit most of the training dataset points and hence gets more complex. To resolve high variance issue we need to work on.

- Getting more training data
- Increase Regularization term
- Modify Model Architecture(Neural network architecture)




Mean squared error

- Let the variable we are trying to predict as Y and other covariates as X . We assume there is a relationship between the two such that:
 - $Y = f(X) + e$

- Assume a model $\hat{f}(X)$ of $f(X)$
- Expected squared error at a point x is

$$Err(x) = E \left[(Y - \hat{f}(x))^2 \right]$$



Some experiments with noisy data T1 and T2

Bias-Variance In Terms of MSE

- The $Err(x)$ can be further decomposed as

$$Err(x) = \left(E[\hat{f}(x)] - f(x)\right)^2 + E\left[\left(\hat{f}(x) - E[\hat{f}(x)]\right)^2\right] + \sigma_e^2$$

$$Err(x) = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$

Bias-Variance Tradeoff

- **Overfitting** gives too much predictive power even to noise elements
- Attempt to reduce overfitting can also begin to **underfit**

Bias-Variance Tradeoff (Cont.)

- Low Bias and Low Variance
 - Perfect model
- Low Bias and High Variance
 - Inconsistent models
- High Bias and Low Variance
 - Consistent but inaccurate models
- High Bias and High Variance
 - Inaccurate and inconsistent models





Fig. 1 Graphical illustration of bias and variance.


Bias-Variance Tradeoff

In terms of model complexity

- For the case of **high bias**, a linear model is used.
- And for the case of **high variance**, the model used was super complex.



© Machine Learning @ Berkeley




© Machine Learning @ Berkeley

Bias-Variance Tradeoff

In terms of model complexity

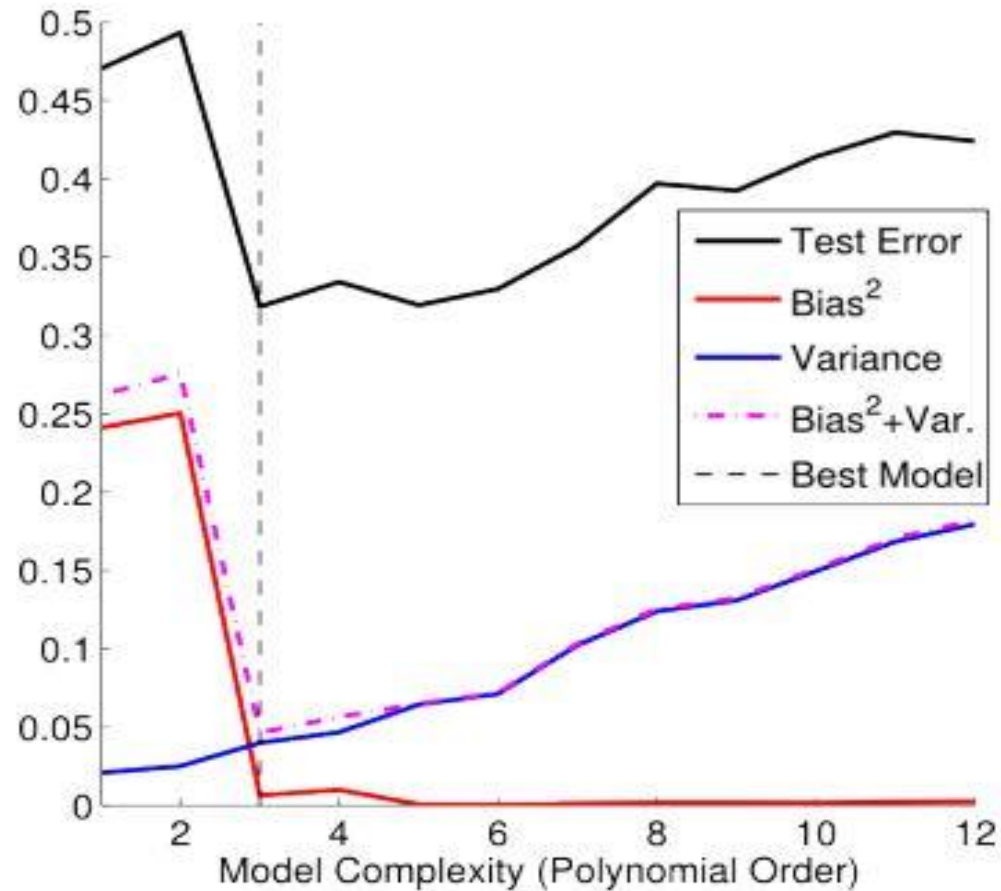
- Low complexity model- Will be prone to underfitting because of high bias and low variance
- High complexity model(Decision trees)- Will be prone to overfitting due to low bias and high variance



© Machine Learning @ Berkeley

Bias-Variance Tradeoff

In terms of model complexity



Balance between Bias-Variance

Regularization is one way to control Bias and Variance

- Which reduces the complexity in the model either by getting rid of the complex features or reducing their importance


Impact of Regularization

For example if the price of a house is based on 4 features which are Location (X1), Number of bedrooms(X2), Year of Construction(X3), Nearby School ranking(X4).

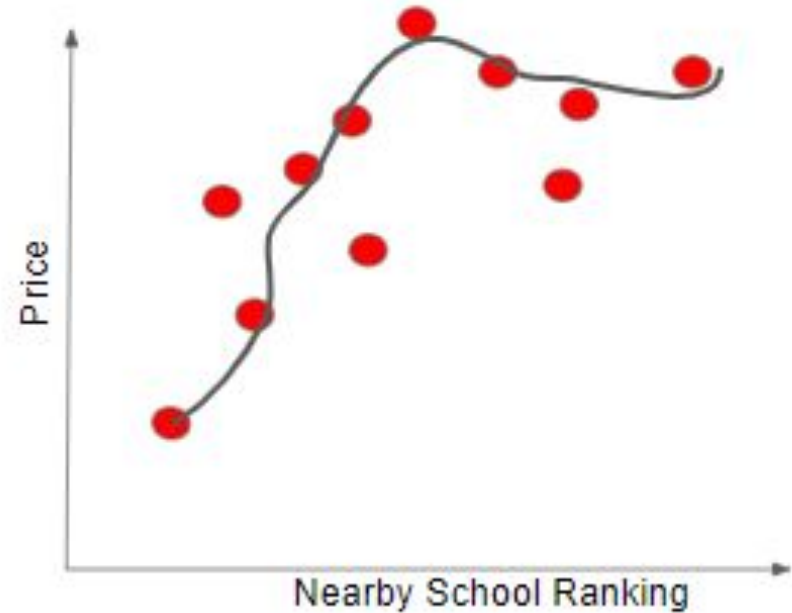
$$Y = 2.5 * X_1 + 3 * X_1 * X_2 + 1.4 * X_3^2 + 4.5 * X_4^3 + 1.3$$

- It reduces the importance of features especially features such as X3 and X4.

Impact of Regularization(Cont.)



Without Regularization



With Regularization

Regularization process

Optimization objective of Linear Regression.

$$W^* = \operatorname{argmin} \left(\frac{1}{2n} * \left(\sum_{i=1}^n (f(X_i) - Y_i)^2 + \lambda \sum_{j=1}^m (W_j)^2 \right) \right)$$

Regularization Process

Optimization objective of Linear Regression.


$$W^* = \operatorname{argmin} \left(\frac{1}{2n} * \left(\sum_{i=1}^n (f(X_i) - Y_i)^2 + \lambda \sum_{j=1}^m (W_j)^2 \right) \right)$$

- Thus, λ acts as a hyperparameter to control the Bias- Variance trade-off.

Regularization Techniques


- **L1 / Lasso Regression**
 - adds absolute value of weights

$$= \sum_{i=1}^N \left\{ y_i - \sum_{j=0}^M w_j x_{ij} \right\}^2 + \lambda \sum_{j=0}^M |w_j|$$



- **L2/ Ridge Regression**
 - Squared value of weights

$$\lambda \sum_{j=1}^m (w_j)^2$$




Regularization Techniques

- **Elastic Net**

- Elastic Net includes both L1 and L2 norm regularization terms.

$$\hat{\beta} \equiv \underset{\beta}{\operatorname{argmin}} (\|y - X\beta\|^2 + \lambda_2 \|\beta\|^2 + \lambda_1 \|\beta\|)$$



Other Techniques to control Bias-Variance Tradeoff

- Feature selection
- Randomization
- Increase data
- Early stopping
- Choice of Algorithm

From Trees to Forests


Random Trees:

Iteratively partition the data and minimize the RSS of the partitions (for regression)

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

Key points:

- For every region we use the mean or mode to compute the prediction value.
- Top to bottom greedy algorithm that might not yield the best tree.



From Trees to Forests cont'd

Why do we stop?

- We could have as many partitions as to fit single observations in the data, basically memorize the training data -> Overfit

When do we stop?

- Naively set a minimum number of samples per partition
- Naively set a minimum RSS improvement per iteration

Better

Grow a large tree and prune it back -> Regularization

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

From Trees to Forests cont'd

Even better : use an ensemble

- Bagging
- Random Forests
- Boosting

Combine “weak learners” with high bias into a model that has lower bias (Boosting)

Combine “strong learners” with low bias (high variance) into a general model with lower variance (Bagging)

Random forests are somewhere in the middle.

What's the price?

- Mainly interpretability and more computation

Neural Networks

Overfitting machines

- Large number of parameters
- Variable architectures
- They can theoretically approximate any function (Cybenko, Hornik)

Deep learning just makes things even harder

- Deep architectures
- Millions(Billions) of parameters

Lu et.al proved in 2017 expressivity on Lebesgue integrable functions for width limited deep architectures, still the depth is a variable.

Neural Networks

Strategies

- Naively opt for simpler architectures
- Early stopping - stop when the validation error starts to increase
- L2 regularization - Penalize the parameters

$$R(\theta) = \sum_{i \in Tr} (y_i - f(\mathbf{x}_i; \theta))^2 + \lambda \sum_{j=1}^p \theta_j^2$$

- Ensembles
- Dropout
- Data augmentation
- SGD
- Bayesian networks

References

- gpu image: https://miro.medium.com/max/3322/1*WBMLfNKon41AFt8lpwbFFw.png
- data image: <https://wallpaperbro.com/img/451165.jpg>
- data exp:
https://www.ft.com/__origami/service/image/v2/images/raw/https%3A%2F%2Fs3-eu-west-1.amazonaws.com%2Fic-ez-prod%2Fez%2Fimages%2F9%2F8%2F1%2F3%2F3173189-1-eng-GB%2FSectorFocus_171117_line.jpg?source=invchron
- Zettabyte: <https://en.wikipedia.org/wiki/Zettabyte>
- curse: https://miro.medium.com/proxy/1*kpNRfZ9DgkNPf_dA6DukLw.png
- curse acc: https://miro.medium.com/max/535/1*vah8loINqlxNHq9ysVzYkw.png
- lin reg complexity: <https://www.thekerneltrip.com/machine/learning/computational-complexity-learning-algorithms/>
- complex: https://miro.medium.com/max/1663/1*0BvHowg6TmamenKHVsxnGA.png
- flip h: <https://i.stack.imgur.com/wM1Kc.png>
- <https://github.com/ajaymache/machine-learning-yearning/tree/master/full%20book>
- Irreducible error: <https://heartbeat.fritz.ai/bias-variance-tradeoff-to-avoid-under-overfitting-d92f1fcff352>
- Bias/var: <https://medium.com/ml-research-lab/chapter-3-bias-and-variance-trade-off-in-machine-learning-a449fa1e2729>
- Bias/var: <https://medium.com/datadriveninvestor/bias-variance-trade-off-fb5fa4c8ab56>
- variance: <https://www.knowledgehut.com/blog/data-science/bias-variance-tradeoff-in-machine-learning>
- Avoidable bias: <http://upscfever.com/upsc-fever/en/data/deeplearning3/8.html>
- <https://harangdev.github.io/deep-learning/structuring-machine-learning-projects/16/>
- Under/Overfitting: <https://medium.com/greyatom/what-is-underfitting-and-overfitting-in-machine-learning-and-how-to-deal-with-it-6803a989c76>
- bias/var image: <https://cscherrer.github.io/post/bias-variance/>
- bias/var tradeoff image: <http://scott.fortmann-roe.com/docs/BiasVariance.html>
- bias/var regression image: <https://towardsdatascience.com/holy-grail-for-bias-variance-tradeoff-overfitting-underfitting-7fad64ab5d76>
- bias/var complexity: <https://insidebigdata.com/2014/10/22/ask-data-scientist-bias-vs-variance-tradeoff/>
- regularization image1: <https://medium.com/swlh/the-what-why-and-how-of-bias-variance-trade-off-55c0b6cf3d00>
- regularization image2: Timothy Downey, Lecture notes ML2 WiSe 2018, Beuth University of Applied Sciences.
- regularization image3: https://web.stanford.edu/~hastie/TALKS/enet_talk.pdf
- stop learning: http://fouryears.eu/wp-content/uploads/2017/12/early_stopping.png
- Timothy Downey, Lecture notes ML2 WiSe 2018, Beuth University of Applied Sciences.
- A.Kryzhevsky et. al ImageNet classification with Deep Convolutional Neural Networks
- James et. al An introduction to statistical learning with applications in R